Predicting Earthquake Ground-Motion Directionality Differences between Nearby Sites using Machine Learning

UE24CS352A – Machine Learning · Mini-Project
Deeksha G (PES1UG24AM078) | **Harshitha B** (PES1UG24AM118)
Code: https://github.com//earthquake-directionality-ml

1. Problem Statement

During an earthquake, shaking at a site is strongest along one horizontal orientation (RotD100). Building codes treat this orientation as uniformly random, but observations show that nearby sites have correlated orientations. Our task (assigned problem, based on Poulos & Silva-Lopez) is to build a model *f* that takes the seismological, geological and geographical attributes of two sites and of the earthquake, and predicts the absolute difference $\Delta\theta$ (0–90°) between their orientations of maximum shaking. We compare a distance-only benchmark, ridge regression and a neural network.

2. Dataset

We use the NGA-West2 database (PEER): a flatfile with 274 attributes per record plus the two-component acceleration time series of shallow crustal earthquakes worldwide. Records from earthquakes with magnitude < 5.0 were removed (____ records kept). $\Delta\theta$ is not in the flatfile, so we computed the RotD100 orientation of each record from its waveforms and paired all stations that recorded the same earthquake and are within 5 km of each other, giving ____ examples from ____ earthquakes. Features include earthquake source data (magnitude, mechanism, rupture size, depth), distance measures, site properties (V_{s30}, NEHRP class), ground-motion intensities (PGA, PGV, PGD) and the distance between the two sites. Numeric features were standardised (zero mean, unit variance) and categorical ones one-hot encoded. Split: 80% train, 10% validation, 10% test (____ / ____ / ____).

3. Approach

Benchmark: linear regression $\Delta\theta = c_1 + c_2 d$ using only the inter-site distance. **Ridge regression:** least squares with an L2 penalty on the coefficients; features were chosen by a backward analysis (rank each feature by the drop in validation R^2 when it is removed) followed by a forward analysis (add features in that order until validation R^2 stops improving), alternating with tuning of the regularisation parameter α on the validation set. **Neural network:** a multilayer perceptron (ReLU, Adam) using all features, with a grid search over number of layers, neurons per layer and L2 strength, selected by validation R^2. Models are compared by R^2 and mean absolute error (MAE) on the test set.

4. Implementation Overview

Python with NumPy, pandas, SciPy and scikit-learn. *waveform.py* reads the .AT2 files and finds the RotD100 orientation by passing both horizontal components through a damped single-degree-of-freedom oscillator and rotating the response through 0–179°; *pairs.py* builds the site pairs and $\Delta\theta$; *modeling.py* contains the split, preprocessing, feature selection, ridge model and neural network; *run_pipeline.py* runs everything and saves metrics, figures and models; *demo.py* predicts $\Delta\theta$ for two chosen sites. Setup and run steps are in the repository README.

5. Results

Model	R^2 (train)	R^2 (test)	MAE (test, deg)
Benchmark	____	____	____
Ridge regression	____	____	____
Neural network	____	____	____

6. Conclusions

[Fill after running the pipeline: which model performed best, whether the models capture the mean trend of $\Delta\theta$ with distance and PGA, overfitting, remaining uncertainty and future work.]

7. Contributions

[Describe who did what, e.g. Deeksha: ____ ; Harshitha: ____ .]

Reference: Poulos, A. and Silva-Lopez, R., Detecting and predicting earthquake ground motion directionality patterns using machine learning tools (Stanford); Ancheta et al. (2014), NGA-West2 database, Earthquake Spectra.